COSC3360 - Fundamentals of Operating Systems

Programming Assignment 1

For this assignment, you will create a parallel fixed-length code decompressor using the tools we learned in class to create multiple processes and threads.

Computers represent a character as a sequence of 8 bits, which means you can represent up to 256 symbols per character. However, depending on the source, it is common to have messages/files using an alphabet with a size of fewer than 256 symbols.

The simplest way to achieve compression is to use a fixed-length code generator. The idea behind this type of code generator is to assign a fixed-length bit sequence to each alphabet symbol.

Given an alphabet (where each symbol is represented by a character and a decimal value with the code of the symbol), you need to implement a fixed-length code decompressor based on the following steps:

- Read the contents of the input file using input redirection (STDIN). The format of the input file is as follow:

    - An integer value representing the number of symbols in the alphabet

    - n lines (where n is the number of symbols in the alphabet) with a char representing the value of the symbol and an integer value representing the symbol's code (in decimal notation).

    - A string representing the compressed message (sequence of bits).

        Given the previous format, the following file represents a valid input file:

        3
        a 2
        b 4
        c 5
        010010010100101101101

- Calculate the number of bits of the fixed-length codes based on the value of the greatest base 10 code in the alphabet. For decimal values greater than zero, you can use the following formula to calculate the number of bits of the fixed-length codes: ceiling(log_2(greatest_base_10_code_in_the_alphabet +1)).

- Create n child processes or threads (where n is the number of symbols in the alphabet). Each child process or thread executes the following tasks:

  - Determines the binary representation of the symbol's code.

  - Determines the frequency of the symbol in the compressed message.

- Create m child processes or threads (where m is the number of characters in the decompressed message) to determine each character of the decompressed message.
- Print the information about the alphabet and the decompressed message. Given the previous input file, the expected output is:

Alphabet:
Character: a, Code: 010, Frequency: 3
Character: b, Code: 100, Frequency: 1
Character: c, Code: 101, Frequency: 3

Decompressed message: aaabccc

NOTES:

- You must select the appropriate tool to implement the parallel fixed-length code decompressor. Not using multiple POSIX threads or multiple processes will translate into a penalty of 100%.
- You must use the output statement format based on the example above. Only printable characters will be considered valid input to guarantee consistency when printing the alphabet's characters to STDOUT.
- You can safely assume that the input will always be valid.